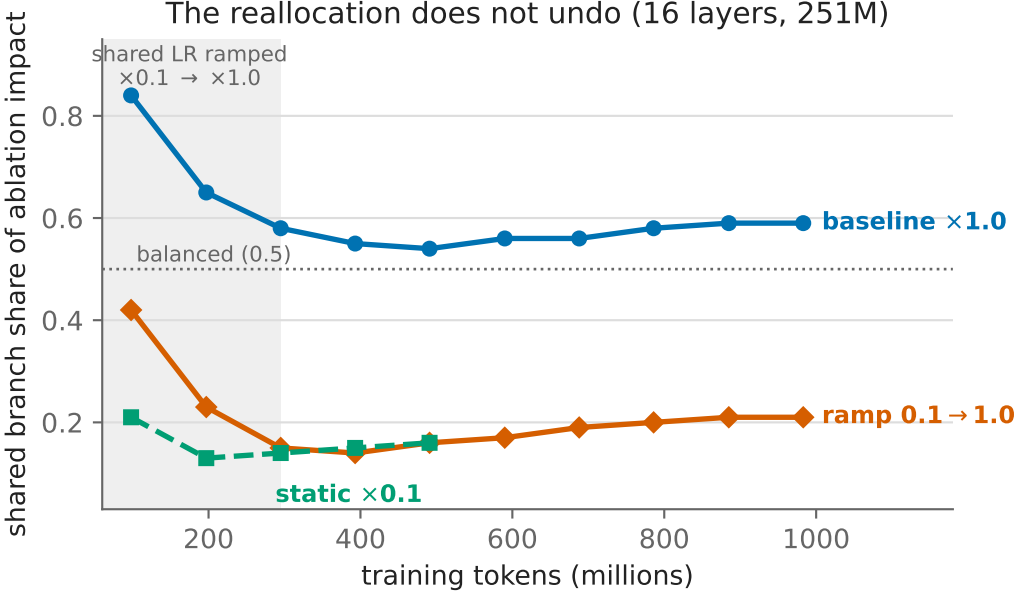
The reallocation does not undo (16 layers, 251M)
shared LR ramped
×0.1 → ×1.0
shared branch share of ablation impact
0.8
0.6
0.4
0.2
balanced (0.5)
baseline ×1.0
ramp 0.1→1.0
static ×0.1
200
400
600
800
1000
training tokens (millions)